

SAATHI

Your Family's Health Memory

Target	Core User	Phase 0 Cost	Phase 0 Timeline
Urban India, 25–35	Wearable owner + family caregiver	~₹2,200 / month	6-8 weeks

1. Problem Statement

Health data - from wearables, lab reports, prescriptions, and outpatient records is fragmented across chats, clinic folders, and siloed apps. This scattering makes it difficult to maintain a complete patient history and ask critical follow-up questions, especially when the caregiver managing this context is remotely located.

The core issue is not just missing health data, but the lack of an intelligence layer to unify data sources into a coherent picture, accessible to the right person, in their language, precisely when needed.

2. Product Vision

Saathi is a shared family health memory. It reads your reports, tracks your numbers, and ensures that you're prepared when a health moment arrives - an appointment, a new diagnosis, a confusing lab report, or a body signal; so you and your family have the full picture, in your language, without scrambling.


2.1 The Trifecta: Three Data Sources, One Intelligence Layer

The product is built on the convergence of three inputs, that currently live in **isolation**:

Health Dashboard	Lab Reports	Doctor Visits
Wearable devices, tracking Steps, Sleep, HR is auto-ingested daily.	Uploaded PDFs are compared to previous and wearables data.	Doctor prescriptions with notes help build the patient's record.

2.2 The Retention Loop

The product has a natural cyclical mechanic that does not require **manufactured habit formation**:



Each stage motivates the next, together mirror how health reports should work cyclically -

- The Health Digest keeps the context populated between appointments.
- The pre-appointment brief makes that context valuable and on-hand at the appointment.

2.3 User Journeys

Stage	Journey	What Happens	Actionable Outcome
Acquisition	The Pre-Appointment Brief	Parent's blood panel is sent to Saathi via WhatsApp. Saathi parses, cross-references it with past reports, and flags new concerns for the child.	Exportable one-page brief with informed questions, shareable via WhatsApp to patient.
Activation	Document Clarity (WhatsApp)	English 4-page hospital discharge summary sent to Saathi. Parent gets a regional summary: details, medications, risks, and follow-up.	Plain-language summary for secondary user. Structured version auto-shared with primary.
Retention	Weekly / Monthly Health Digest	Monday's resting HR was stable, but sleep efficiency and mid-week step count dropped. This low energy aligns with the previous lab report showing borderline Vitamin D.	Retest suggestion + one lifestyle nudge. Keeps the data pipeline active between appointments.
Retention	Post-Diagnosis Tracking (Chronic Illness)	The latest HbA1c is 6.8 (pre-diabetic). Over 12 weeks, a digest will track sleep, activity, and weight, providing trend data for the 3-month follow-up.	12-week trend summary follow-up brief, with actionable lifestyle changes, suggestions.
Engagement	Post-Appointment Capture	New prescription issued after visit. User uploads it to Saathi - medication, follow-up date, doctor's instructions stored. Feeds next digest and next brief automatically.	Medication reminder set, follow-up flagged, next brief pre-loaded. Closes the data pipeline loop.
Engagement	Talk to Your Data	User asks: "Has my sleep improved since I started 8,000 steps?" Saathi pulls wearable history, cross-references last blood test, surfaces one grounded, cited answer. Never interprets medically.	A specific question or observation to bring to the next appointment. Converts passive data into active health agency.

3. Users & Privacy Architecture

3.1 Primary User

- **Demographics:** Urban Indians (25-35) interested in or owning wearables and using AI (e.g., ChatGPT, Claude), often for remote family health management.
- **Motivation:** Reduce the cognitive load and anxiety of being the family's health coordinator from afar.
- **Acquisition entry point:** A specific upcoming appointment or confusing report is discovered via Google, word-of-mouth, or hospital/clinic partnerships.
- **Wearable as filter:** Wearable owners already value tracking; the app intelligently enhances their behavior. Non-wearable users are prompted to connect Fitbit or obtain a basic device.

3.2 Secondary User


- **Who they are:** The family member being cared for - typically a parent, 50+, may not be tech-savvy.

- **Interaction model:** Simple, voice-first, regional language App and WhatsApp to send a document and get a plain-language summary.
- **Privacy architecture:** The secondary user consents at onboarding to share a health summary report with the primary user via the app or WhatsApp. This report is generated, not automatically synced, and is ingested by the primary user into their Saathi context.
- **DPDP Act 2023:** Health data is sensitive personal data stored on AWS ap-south-1 (Indian servers). A consent management layer is required pre-launch. Legal review is advised for the Sharing feature.

4. System Architecture

4.1 Design Principles

- **Context-first:** LLM Responses use only retrieved personal data; medical advice is blocked.
- **Passive ingestion:** The system gathers context automatically. Wearables sync post-approval. The secondary user only needs to forward WhatsApp messages.
- **Language-native:** Input is accepted in any supported language, processed internally in English, and output is back-translated.



4.2 Document Parsing - Phase Priority

Phase	Document Type	Parser	Guardrail
Phase 0	Lab reports (Max, Apollo, Thyrocare, Dr. Lal)	LlamaParse + structured extraction	Trend display only - avoid interpretation
Phase 0	Wearable data (Google Fit, Fitbit)	Native API	None needed, metrics only
Phase 1	Prescriptions (digital, typed)	LlamaParse	Medication list only (no dosage advice)
Phase 1	Hospital discharge summaries	LlamaParse + Bolna cross-lingual RAG	Plain language summary, no clinical interpretation
Phase 2	Insurance documents, billing	Bolna cross-lingual RAG or similar	Factual extraction only
Out of scope	Handwritten prescriptions / older reports	Not reliable with current OCR	Exclude from Phase 0–1 scope

5. Tech Stack

Layer	Tool	Why	Phase 0 Alternative
LLM	Claude Haiku (chat) + Claude Sonnet (briefs)	Best instruction-following; Haiku cost-effective at scale	OpenAI small and reasoning models
Orchestration	LangGraph (Python)	Stateful agent graphs; multi-step RAG + guard flows native	LangChain basic chains
Speech ASR	Sarvam AI - Saarika	Best Indic-native ASR in production; free tier	Whisper large-v3 or Elevenlabs
Translation	IndicTrans2 (IITM, open-source)	22 Indian languages; self-hostable; zero marginal cost	Sarvam Translate API
Cross-lingual RAG	IndicTrans2 + pgvector embeddings	Validated by Bolna YC F25 approach for Indian docs	Bolna API (if B2B pivot)
Vector + DB + Auth	Supabase (pgvector)	Free tier; Postgres-native; auth + vector in one	Chroma (local) + Firebase
OCR / Doc Parse	LlamaParse + Tesseract	LlamaParse for structured PDFs; Tesseract for photo fallback	Google Document AI (higher cost)
Wearable Data	Google Fitbit API	Free; covers Android-majority Indian market	Garmin (needs domain)
Backend	FastAPI (Python)	Async; native LangChain ecosystem; simple deploy	–
Frontend	Next.js 15 (Mweb app)	Mweb-first; offline capable; single codebase web + mobile	–
WhatsApp Ingestion	WhatsApp Business API (Meta)	Native share-sheet integration; parent forwards PDF naturally	Manual upload fallback
Monitoring	LangSmith	LLM call tracing; hallucination audit; free tier	Langfuse (open-source)
Deployment Phase 0	Vercel	Free; GitHub deploy; perfect for solo dev	Railway (\$5/month)
Deployment Phase 1+	AWS ECS + RDS ap-south-1	Indian data residency for DPDP compliance; cost scales well	GCP Cloud Run

6. Cost Analysis

Phase 0-1 = 50 active users - 500 closed beta

API costs assume 10 messages/user/day, ~600 tokens per exchange, plus weekly digest generation.

Cost Item	Phase 0(₹/mo)	Phase 1(₹/mo)
Claude Haiku: daily chat queries	~₹1,400	~₹8,000
Claude Sonnet: weekly digests + briefs	~₹400	~₹3,000
Sarvam AI ASR: voice input	Free	~₹1,200

IndicTrans2: self-hosted Railway	~₹400	~₹2,000
Supabase: DB + vector + auth	Free	~₹1,700
Vercel / AWS compute	Free	~₹3,400
LlamaParse: lab report OCR	Free	~₹700
WhatsApp Business API	Free	~₹500
LangSmith: monitoring	Free	Free
TOTAL ESTIMATED	~₹2,200	~₹20,500

6.1 Revenue Model

- **Free:** Core health dashboard + 2 family + 3 reports per month + Monthly digest, pre-appointment brief export.
- **Individual Plan ₹89/month:** Free + 3 family + 5 reports/mo, Weekly digest, Voice input. Break-even: ~230 paying users at Phase 1.
- **Family Plan ₹179/month (Phase 2):** Individual plan + 5 family, shared health graph, WhatsApp ingestion for secondary users, monthly family health summary.
- **Wearable partnership (Phase 1):** Co-marketing with Fitbit / Garmin India - Saathi as the intelligence layer for their hardware. Distribution channel and marketing uplift.

7. Build Timeline

Phase	Period	Deliverables
Phase 0 MVP	Weeks 1–8	Next.js web-app + FastAPI backend · Supabase setup · LLM integration · Lab report upload + LlamaParse · Fitbit/other wearable sync · RAG layer across wearable + lab data · Weekly digest (text) · Pre-appointment brief export · English + Hindi text · WhatsApp Business API ingestion · LangSmith monitoring
Phase 1 Beta	Months 3–5	Sarvam ASR voice input · IndicTrans2 translation layer · Medical guard classifier · Prescription parsing · DPDP legal review · Paid Plans + Razorpay integration
Phase 2 Public	Months 6–9	3 additional languages (Tamil, Telugu, Bengali) · Family plan · Monthly health digest · Wearable onboarding CTA for non-owners · Product Hunt India launch · Discharge summary + insurance parsing

8. Feasibility Assessment

Dimension	Rating	Pre-build Assessment
Core chat + RAG (Phase 0)	HIGH	Well-understood stack. Supabase + LangGraph + LLM to lead.
Wearable data integration	HIGH	Fitbit API is free, well-documented.
Hindi language support	HIGH	Sarvam + IndicTrans2 is solid for Hindi. English-first with Hindi overlay.
Weekly digest + pre-appt brief	MED-HIGH	Most LLMs with defined frontend are good at synthesis.
Post-Appointment Capture	MED-HIGH	Voice log/prescription upload is easy. Reliable, structured extraction (meds, dosages, dates) from written docs is harder.
WhatsApp Business API ingestion	MED-HIGH	Meta's API is well-documented. Approval process will be lengthy.

Talk to Your Data	MED-HIGH	RAG on structured wearable and lab data with strict context-only prompting is technically feasible using the existing stack.
Post-Diagnosis Chronic Tracking	MED-HIGH	The feature needs a min of 8-10 full weeks of consistent data from the synced wearable devices and corresponding report generations to deliver its intended, complete value to the user.
Medical guard classifier	MEDIUM	Achievable but requires deliberate building. Cannot be bolted on at the end.
Handwritten / old report OCR	LOW	Explicitly out of scope. Current OCR accuracy on degraded scans will produce wrong health data. Worse than not supporting it.
Execution through Phase 1	MEDIUM	Phase 0 is achievable with small team, with AI-assisted development. Phase 1 requires some support.
DPDP Act compliance	MEDIUM	Health = sensitive personal data. Indian server residency is achievable on AWS ap-south-1. Legal review will be regulatory for Phase 1.
Monetisation at ₹99/month	MEDIUM	Indian willingness to pay for health AI is unproven at scale. Free plan with a clear upgrade trigger is essential. This is the biggest market risk.

8.1 Market Validation

Bolna (YC F25), building cross-lingual RAG infrastructure for India's document-heavy sectors - insurance, banking, hospital discharge summaries - recently launched multilingual embedding supporting 11 Indian languages. Their core thesis: over 90% of Indians cannot read English, yet every major document they interact with is written in it. This validates both the technical approach and the market urgency. Saathi applies this infrastructure specifically to personal and family health, a domain Bolna is not yet targeting.

9. Conclusion

Saathi addresses the gap between India's rapid urbanisation and the accelerating AI adoption among the 25-35 demographic. It leverages existing user moments of health anxiety and confusion, rather than relying on behaviour change. With a proven, cost-efficient technical foundation and a specific go-to-market strategy, Saathi acquires users with intent. The product is capital-light, with a Phase 0 under ₹3,000/month and a path to a ₹179/month family plan that breaks even at early adoption by users. Its architecture is defensible, becoming more irreplaceable with use. The opportunity is to own the personal health memory layer for 300 million urban Indians before a general-purpose AI competitor emerges.